

2.10 — Introduction to the preprocessor

 learncpp.com/cpp-tutorial/introduction-to-the-preprocessor/

When you compile your project, you might expect that the compiler compiles each code file exactly as you've written it. This actually isn't the case.

Instead, prior to compilation, each code (.cpp) file goes through a **preprocessing** phase. In this phase, a program called the **preprocessor** makes various changes to the text of the code file. The preprocessor does not actually modify the original code files in any way -- rather, all changes made by the preprocessor happen either temporarily in-memory or using temporary files.

As an aside...

Historically, the preprocessor was a separate program from the compiler, but in modern compilers, the preprocessor may be built right into the compiler itself.

Most of what the preprocessor does is fairly uninteresting. For example, it strips out comments, and ensures each code file ends in a newline. However, the preprocessor does have one very important role: it is what processes **#include** directives (which we'll discuss more in a moment).

When the preprocessor has finished processing a code file, the result is called a **translation unit**. This translation unit is what is then compiled by the compiler.

Related content

The entire process of preprocessing, compiling, and linking is called **translation**.

If you're curious, here is a list of [translation phases](#). As of the time of writing, preprocessing encompasses phases 1 through 4, and compilation is phases 5 through 7.

Preprocessor directives

When the preprocessor runs, it scans through the code file (from top to bottom), looking for preprocessor directives. **Preprocessor directives** (often just called *directives*) are instructions that start with a # symbol and end with a newline (NOT a semicolon). These directives tell the preprocessor to perform certain text manipulation tasks. Note that the preprocessor does not understand C++ syntax -- instead, the directives have their own syntax (which in some cases resembles C++ syntax, and in other cases, not so much).

In this lesson, we'll look at some of the most common preprocessor directives.

As an aside...

Using directives (introduced in lesson [2.9 -- Naming collisions and an introduction to namespaces](#)) are not preprocessor directives (and thus are not processed by the preprocessor). So while the term **directive** *usually* means a **preprocessor directive**, this is not always the case.

#Include

You've already seen the **#include** directive in action (generally to **#include <iostream>**). When you **#include** a file, the preprocessor replaces the **#include** directive with the contents of the included file. The included contents are then preprocessed (which may result in additional **#includes** being preprocessed recursively), then the rest of the file is preprocessed.

Consider the following program:

```
#include <iostream>

int main()
{
    std::cout << "Hello, world!\n";
    return 0;
}
```

When the preprocessor runs on this program, the preprocessor will replace **#include <iostream>** with the contents of the file named "iostream" and then preprocess the included content and the rest of the file.

Once the preprocessor has finished processing the code file plus all of the **#included** content, the result is called a **translation unit**. The translation unit is what is sent to the compiler to be compiled.

Key insight

A translation unit contains both the processed code from the code file, as well as the processed code from all of the **#included** files.

Since **#include** is almost exclusively used to include header files, we'll discuss **#include** in more detail in the next lesson (when we discuss header files).

Macro defines

The **#define** directive can be used to create a macro. In C++, a **macro** is a rule that defines how input text is converted into replacement output text.

There are two basic types of macros: *object-like macros*, and *function-like macros*.

Function-like macros act like functions, and serve a similar purpose. Their use is generally considered unsafe, and almost anything they can do can be done by a normal function.

Object-like macros can be defined in one of two ways:

```
#define IDENTIFIER
#define IDENTIFIER substitution_text
```

The top definition has no substitution text, whereas the bottom one does. Because these are preprocessor directives (not statements), note that neither form ends with a semicolon.

The identifier for a macro uses the same naming rules as normal identifiers: they can use letters, numbers, and underscores, cannot start with a number, and should not start with an underscore. By convention, macro names are typically all upper-case, separated by underscores.

Object-like macros with substitution text

When the preprocessor encounters this directive, any further occurrence of the identifier is replaced by *substitution_text*. The identifier is traditionally typed in all capital letters, using underscores to represent spaces.

Consider the following program:

```
#include <iostream>

#define MY_NAME "Alex"

int main()
{
    std::cout << "My name is: " << MY_NAME << '\n';

    return 0;
}
```

The preprocessor converts the above into the following:

```
// The contents of iostream are inserted here

int main()
{
    std::cout << "My name is: " << "Alex" << '\n';

    return 0;
}
```

Which, when run, prints the output **My name is: Alex**.

Object-like macros with substitution text were used (in C) as a way to assign names to literals. This is no longer necessary, as better methods are available in C++. Object-like macros with substitution text should generally now only be seen in legacy code.

We recommend avoiding these kinds of macros altogether, as there are better ways to do this kind of thing. We discuss this more in lesson [5.1 -- Constant variables \(named constants\)](#).

Object-like macros without substitution text

Object-like macros can also be defined without substitution text.

For example:

```
#define USE_YEN
```

Macros of this form work like you might expect: any further occurrence of the identifier is removed and replaced by nothing!

This might seem pretty useless, and it *is useless* for doing text substitution. However, that's not what this form of the directive is generally used for. We'll discuss the uses of this form in just a moment.

Unlike object-like macros with substitution text, macros of this form are generally considered acceptable to use.

Conditional compilation

The *conditional compilation* preprocessor directives allow you to specify under what conditions something will or won't compile. There are quite a few different conditional compilation directives, but we'll only cover the three that are used by far the most here: *`#ifdef`*, *`#ifndef`*, and *`#endif`*.

The *`#ifdef`* preprocessor directive allows the preprocessor to check whether an identifier has been previously *`#defined`*. If so, the code between the *`#ifdef`* and matching *`#endif`* is compiled. If not, the code is ignored.

Consider the following program:

```

#include <iostream>

#define PRINT_JOE

int main()
{
#ifdef PRINT_JOE
    std::cout << "Joe\n"; // will be compiled since PRINT_JOE is defined
#endif

#ifdef PRINT_BOB
    std::cout << "Bob\n"; // will be excluded since PRINT_BOB is not defined
#endif

    return 0;
}

```

Because PRINT_JOE has been *#defined*, the line `std::cout << "Joe\n"` will be compiled. Because PRINT_BOB has not been *#defined*, the line `std::cout << "Bob\n"` will be ignored.

#ifndef is the opposite of *#ifdef*, in that it allows you to check whether an identifier has *NOT* been *#defined* yet.

```

#include <iostream>

int main()
{
#ifndef PRINT_BOB
    std::cout << "Bob\n";
#endif

    return 0;
}

```

This program prints “Bob”, because PRINT_BOB was never *#defined*.

In place of *#ifdef PRINT_BOB* and *#ifndef PRINT_BOB*, you’ll also see *#if defined(PRINT_BOB)* and *#if !defined(PRINT_BOB)*. These do the same, but use a slightly more C++-style syntax.

#if 0

One more common use of conditional compilation involves using *#if 0* to exclude a block of code from being compiled (as if it were inside a comment block):

```

#include <iostream>

int main()
{
    std::cout << "Joe\n";

    #if 0 // Don't compile anything starting here
        std::cout << "Bob\n";
        std::cout << "Steve\n";
    #endif // until this point

    return 0;
}

```

The above code only prints “Joe”, because “Bob” and “Steve” are excluded from compilation by the `#if 0` preprocessor directive.

This provides a convenient way to “comment out” code that contains multi-line comments (which can’t be commented out using another multi-line comment due to multi-line comments being non-nestable):

```

#include <iostream>

int main()
{
    std::cout << "Joe\n";

    #if 0 // Don't compile anything starting here
        std::cout << "Bob\n";
        /* Some
         * multi-line
         * comment here
         */
        std::cout << "Steve\n";
    #endif // until this point

    return 0;
}

```

To temporarily re-enable code that has been wrapped in an `#if 0`, you can change the `#if 0` to `#if 1`:

```

#include <iostream>

int main()
{
    std::cout << "Joe\n";

    #if 1 // always true, so the following code will be compiled
        std::cout << "Bob\n";
        /* Some
         * multi-line
         * comment here
         */
        std::cout << "Steve\n";
    #endif

    return 0;
}

```

Object-like macros don't affect other preprocessor directives

Now you might be wondering:

```

#define PRINT_JOE

#ifdef PRINT_JOE
// ...

```

Since we defined *PRINT_JOE* to be nothing, how come the preprocessor didn't replace *PRINT_JOE* in *#ifdef PRINT_JOE* with nothing?

Macros only cause text substitution for non-preprocessor commands. Since *#ifdef PRINT_JOE* is a preprocessor command, text substitution does not apply to *PRINT_JOE* within this command.

As another example:

```

#define F00 9 // Here's a macro substitution

#ifdef F00 // This F00 does not get replaced because it's part of another
preprocessor directive
    std::cout << F00 << '\n'; // This F00 gets replaced with 9 because it's part of
the normal code
#endif

```

However, the final output of the preprocessor contains no directives at all -- they are all resolved/stripped out before compilation, because the compiler wouldn't know what to do with them.

The scope of *#defines*

Directives are resolved before compilation, from top to bottom on a file-by-file basis.

Consider the following program:

```
#include <iostream>

void foo()
{
#define MY_NAME "Alex"
}

int main()
{
    std::cout << "My name is: " << MY_NAME << '\n';

    return 0;
}
```

Even though it looks like `#define MY_NAME "Alex"` is defined inside function `foo`, the preprocessor won't notice, as it doesn't understand C++ concepts like functions. Therefore, this program behaves identically to one where `#define MY_NAME "Alex"` was defined either before or immediately after function `foo`. For readability, you'll generally want to `#define` identifiers outside of functions.

Once the preprocessor has finished, all defined identifiers from that file are discarded. This means that directives are only valid from the point of definition to the end of the file in which they are defined. Directives defined in one code file do not have impact on other code files in the same project.

Consider the following example:

function.cpp:

```
#include <iostream>

void doSomething()
{
#ifdef PRINT
    std::cout << "Printing!\n";
#endif
#ifndef PRINT
    std::cout << "Not printing!\n";
#endif
}
```

main.cpp:


```
void doSomething(); // forward declaration for function doSomething()

#define PRINT

int main()
{
    doSomething();

    return 0;
}
```

The above program will print:

Not printing!

Even though `PRINT` was defined in *main.cpp*, that doesn't have any impact on any of the code in *function.cpp* (`PRINT` is only `#defined` from the point of definition to the end of *main.cpp*). This will be of consequence when we discuss header guards in a future lesson.