

5.8 — Constexpr and consteval functions

 learncpp.com/cpp-tutorial/constexpr-and-consteval-functions/

In lesson [5.5 -- Constexpr variables](#), we introduced the `constexpr` keyword, which we used to create compile-time (symbolic) constants. We also introduced constant expressions, which are expressions that can be evaluated at compile-time rather than runtime.

Consider the following program, which uses two `constexpr` variables:

```
#include <iostream>

int main()
{
    constexpr int x{ 5 };
    constexpr int y{ 6 };

    std::cout << (x > y ? x : y) << " is greater!\n";

    return 0;
}
```

This produces the result:

```
6 is greater!
```

Because `x` and `y` are `constexpr`, the compiler can evaluate the constant expression `(x > y ? x : y)` at compile-time, reducing it to just `6`. Because this expression no longer needs to be evaluated at runtime, our program will run faster.

However, having a non-trivial expression in the middle of our print statement isn't ideal -- it would be better if the expression were a named function. Here's the same example using a function:

```

#include <iostream>

int greater(int x, int y)
{
    return (x > y ? x : y); // here's our expression
}

int main()
{
    constexpr int x{ 5 };
    constexpr int y{ 6 };

    std::cout << greater(x, y) << " is greater!\n"; // will be evaluated at runtime

    return 0;
}

```

This program produces the same output as the prior one. But there's a downside to putting our expression in a function: the call to `greater(x, y)` will execute at runtime. By using a function (which is good for modularity and documentation) we've lost our ability for that code to be evaluated at compile-time (which is bad for performance).

So how might we address this?

Constexpr functions can be evaluated at compile-time

A **constexpr function** is a function whose return value may be computed at compile-time. To make a function a constexpr function, we simply use the `constexpr` keyword in front of the return type. Here's a similar program to the one above, using a constexpr function:

```

#include <iostream>

constexpr int greater(int x, int y) // now a constexpr function
{
    return (x > y ? x : y);
}

int main()
{
    constexpr int x{ 5 };
    constexpr int y{ 6 };

    // We'll explain why we use variable g here later in the lesson
    constexpr int g { greater(x, y) }; // will be evaluated at compile-time

    std::cout << g << " is greater!\n";

    return 0;
}

```

This produces the same output as the prior example, but the function call `greater(x, y)` will be evaluated at compile-time instead of runtime!

When a function call is evaluated at compile-time, the compiler will calculate the return value of the function call at compile-time, and then replace the function call with the return value.

So in our example, the call to `greater(x, y)` will be replaced by the result of the function call, which is the integer value `6`. In other words, the compiler will compile this:

```
#include <iostream>

int main()
{
    constexpr int x{ 5 };
    constexpr int y{ 6 };

    constexpr int g { 6 }; // greater(x, y) evaluated at compile-time and replaced
with return value 6

    std::cout << g << " is greater!\n";

    return 0;
}
```

To be eligible for compile-time evaluation, a function must have a `constexpr` return type, and a call to the function must have arguments that are known at compile time (e.g. are constant expressions).

Author's note

We'll use the term "eligible for compile-time evaluation" later in the article, so remember this definition.

For advanced readers

There are some other lesser encountered criteria as well. These can be found [here](#).

Our `greater()` function definition and function call in the above example meets these requirements, so it is eligible for compile-time evaluation.

`Constexpr` functions can also be evaluated at runtime

Functions with a `constexpr` return value can also be evaluated at runtime, in which case they will return a non-`constexpr` result. For example:

```

#include <iostream>

constexpr int greater(int x, int y)
{
    return (x > y ? x : y);
}

int main()
{
    int x{ 5 }; // not constexpr
    int y{ 6 }; // not constexpr

    std::cout << greater(x, y) << " is greater!\n"; // will be evaluated at runtime

    return 0;
}

```

In this example, because arguments `x` and `y` are not constant expressions, the function cannot be resolved at compile-time. However, the function will still be resolved at runtime, returning the expected value as a non-constexpr `int`.

Key insight

Allowing functions with a constexpr return type to be evaluated at either compile-time or runtime was allowed so that a single function can serve both cases.

Otherwise, you'd need to have separate functions (a function with a constexpr return type, and a function with a non-constexpr return type). This would not only require duplicate code, the two functions would also need to have different names!

So when is a constexpr function evaluated at compile-time?

You might think that a constexpr function would evaluate at compile-time whenever possible, but unfortunately this is not the case.

According to the C++ standard, a constexpr function that is eligible for compile-time evaluation *must* be evaluated at compile-time if the return value is used where a constant expression is required. Otherwise, the compiler is free to evaluate the function at either compile-time or runtime.

Let's examine a few cases to explore this further:

```

#include <iostream>

constexpr int greater(int x, int y)
{
    return (x > y ? x : y);
}

int main()
{
    constexpr int g { greater(5, 6) };           // case 1: always evaluated at
    compile-time
    std::cout << g << " is greater!\n";

    std::cout << greater(5, 6) << " is greater!\n"; // case 2: may be evaluated at
    either runtime or compile-time

    int x{ 5 }; // not constexpr but value is known at compile-time
    std::cout << greater(x, 6) << " is greater!\n"; // case 3: likely evaluated at
    runtime

    std::cin >> x;
    std::cout << greater(x, 6) << " is greater!\n"; // case 4: always evaluated at
    runtime

    return 0;
}

```

In case 1, we're calling `greater()` with constant expression arguments, so it is eligible to be evaluated at compile-time. The initializer of constexpr variable `g` must be a constant expression, so the return value is used where a constant expression is required. Thus, `greater()` must be evaluated at compile-time.

In case 2, the `greater()` function is again being called with constant expression arguments, so it is eligible for compile-time evaluation. However, the return value is not being used in a context that requires a constant expression (operator<< always executes at runtime), so the compiler is free to choose whether this call to `greater()` will be evaluated at compile-time or runtime!

In case 3, we're calling `greater()` with one argument that is not a constant expression. However, this argument has a value that is known at compile-time. Under the as-if rule, the compiler could decide to treat `x` as constexpr, and call evaluate this call to `greater()` at compile-time. But more likely, it will evaluate it at runtime.

In case 4, the value of argument `x` can't be known at compile-time, so this call to `greater()` will always evaluate at runtime.

Note that your compiler's optimization level setting may have an impact on whether it decides to evaluate a function at compile-time or runtime. This also means that your compiler may make different choices for debug vs. release builds (as debug builds typically have optimizations turned off).

Key insight

A constexpr function must be evaluated at compile-time if the return value is used where a constant expression is required. Otherwise, compile-time evaluation is not guaranteed.

Thus, a constexpr function is better thought of as “can be used in a constant expression”, not “will be evaluated at compile-time”.

When a constexpr (or consteval) function is being evaluated at compile-time, any other functions it calls are required to be evaluated at compile-time (otherwise the initial function would not be able to return a result at compile-time).

Key insight

Put another way, we can categorize the likelihood that a constexpr function will actually be evaluated at compile-time as follows:

Always (required by the standard):

- Constexpr function is called where constant expression is required.
- Constexpr function is called from other function being evaluated at compile-time.

Probably (there's little reason not to):

Constexpr function is called where constant expression isn't required, all arguments are constant expressions.

Possibly (if optimized under the as-if rule):

Constexpr function is called where constant expression isn't required, some arguments are not constant expressions but their values are known at compile-time.

Never (not possible):

Constexpr function is called where constant expression isn't required, some arguments have values that are not known at compile-time.

Determining if a constexpr function call is evaluating at compile-time or runtime

C++ does not currently provide any reliable mechanisms to do this.

What about `std::is_constant_evaluated` or `if constexpr`?

Introduced in C++20, `std::is_constant_evaluated()` (defined in the `<type_traits>` header) was intended to allow differentiation of behavior when a function is evaluating at runtime vs compile-time, so you could do something like this:

```
#include <type_traits> // for std::is_constant_evaluated()

constexpr int someFunction()
{
    if (std::is_constant_evaluated()) // if evaluating at compile-time
        // do something
    else // runtime evaluation
        // do something else
}
```

However, as [the paper that proposed this feature](#) indicates, the standard doesn't actually make a distinction between "compile time" and "run time". So instead, `std::is_constant_evaluated()` returns a `bool` indicating whether the current function is executing in a constant-evaluated context. A **constant-evaluated context** (also called a **constant context**) is defined as one in which a constant expression is required (such as the initialization of a `constexpr` variable). In cases where the compiler is required to evaluate a constant expression at compile-time `std::is_constant_evaluated` will `true` as expected.

However, the compiler may also choose to evaluate a `constexpr` function at compile-time in a context that does not require a constant expression. In such cases, `std::is_constant_evaluated()` will return `false` even though the function did evaluate at compile-time. So `std::is_constant_evaluated()` really means "the compiler is being forced to evaluate this at compile-time", not "this is evaluating at compile-time".

Introduced in C++23, `if constexpr` is shorthand for `if (std::is_constant_evaluated())` (that also does not require the `<type_traits>` header). It has the same issue.

Forcing a `constexpr` function to be evaluated at compile-time

There is no way to tell the compiler that a `constexpr` function should prefer to evaluate at compile-time whenever it can (e.g. in cases where the return value is used in a non-constant expression).

However, we can force a `constexpr` function that is eligible to be evaluated at compile-time to actually evaluate at compile-time by ensuring the return value is used where a constant expression is required. This needs to be done on a per-call basis.

The most common way to do this is to use the return value to initialize a `constexpr` variable (this is why we've been using variable 'g' in prior examples). Unfortunately, this requires introducing a new variable into our program just to ensure compile-time evaluation, which is ugly and reduces code readability.

For advanced readers

There are several hacky ways that people have tried to work around the problem of having to introduce a new constexpr variable each time we want to force compile-time evaluation. See [here](#) and [here](#).

However, in C++20, there is a better workaround to this issue, which we'll present in a moment.

Consteval C++20

C++20 introduces the keyword **constexpr**, which is used to indicate that a function *must* evaluate at compile-time, otherwise a compile error will result. Such functions are called **immediate functions**.

```
#include <iostream>

constexpr int greater(int x, int y) // function is now constexpr
{
    return (x > y ? x : y);
}

int main()
{
    constexpr int g { greater(5, 6) };           // ok: will evaluate at compile-
time
    std::cout << g << '\n';

    std::cout << greater(5, 6) << " is greater!\n"; // ok: will evaluate at compile-
time

    int x{ 5 }; // not constexpr
    std::cout << greater(x, 6) << " is greater!\n"; // error: constexpr functions
must evaluate at compile-time

    return 0;
}
```

In the above example, the first two calls to `greater()` will evaluate at compile-time. The call to `greater(x, 6)` cannot be evaluated at compile-time, so a compile error will result.

Best practice

Use `constexpr` if you have a function that must run at compile-time for some reason (e.g. performance).

Using `constexpr` to make constexpr execute at compile-time C++20

The downside of consteval functions is that such functions can't evaluate at runtime, making them less flexible than constexpr functions, which can do either. Therefore, it would still be useful to have a convenient way to force constexpr functions to evaluate at compile-time (even when the return value is being used where a constant expression is not required), so that we could have compile-time evaluation when possible, and runtime evaluation when we can't.

Consteval functions provides a way to make this happen, using a neat helper function:

```
#include <iostream>

// Uses abbreviated function template (C++20) and `auto` return type to make this
// function work with any type of value
// See 'related content' box below for more info (you don't need to know how these
// work to use this function)
constexpr auto compileTime(auto value)
{
    return value;
}

constexpr int greater(int x, int y) // function is constexpr
{
    return (x > y ? x : y);
}

int main()
{
    std::cout << greater(5, 6) << '\n';           // may or may not execute at
    compile-time
    std::cout << compileTime(greater(5, 6)) << '\n'; // will execute at compile-time

    int x { 5 };
    std::cout << greater(x, 6) << '\n';           // we can still call the
    constexpr version at runtime if we wish

    return 0;
}
```

This works because consteval functions require constant expressions as arguments -- therefore, if we use the return value of a constexpr function as an argument to a consteval function, the constexpr function must be evaluated at compile-time! The consteval function just returns this argument as its own return value, so the caller can still use it.

Note that the consteval function returns by value. While this might be inefficient to do at runtime (if the value was some type that is expensive to copy, e.g. `std::string`), in a compile-time context, it doesn't matter because the entire call to the consteval function will simply be replaced with the calculated return value.

Related content

We cover `auto` return types in lesson [10.9 -- Type deduction for functions](#).

We cover abbreviated function templates (`auto` parameters) in lesson [11.8 -- Function templates with multiple template types](#).

Constexpr/constexpr functions are implicitly inline

Because constexpr functions may be evaluated at compile-time, the compiler must be able to see the full definition of the constexpr function at all points where the function is called. A forward declaration will not suffice, even if the actual function definition appears later in the same compilation unit.

This means that a constexpr function called in multiple files needs to have its definition included into each such file -- which would normally be a violation of the one-definition rule. To avoid such problems, constexpr functions are implicitly inline, which makes them exempt from the one-definition rule.

As a result, constexpr functions are often defined in header files, so they can be `#included` into any `.cpp` file that requires the full definition.

Consteval functions are also implicitly inline (presumably for consistency).

Rule

The compiler must be able to see the full definition of a constexpr or consteval function, not just a forward declaration.

Best practice

Constexpr/constexpr functions used in a single source file (`.cpp`) can be defined in the source file above where they are used.

Constexpr/constexpr functions used in multiple source files should be defined in a header file so they can be included into each source file.

Constexpr/constexpr function parameters are not constexpr

The parameters of a constexpr function are not implicitly constexpr, nor may they be declared as `constexpr`.

Key insight

A constexpr function parameter would imply the function could only be called with a constexpr argument. But this is not the case -- constexpr functions can be called with non-constexpr arguments when the function is evaluated at runtime.

Because such parameters are not constexpr, they cannot be used in constant expressions within the function.

```
constexpr int goo(int c)    // c is not constexpr, and cannot be used in constant
expressions
{
    return c;
}

constexpr int foo(int b)    // b is not constexpr, and cannot be used in constant
expressions
{
    constexpr int b2 { b }; // compile error: constexpr variable requires constant
expression initializer

    return goo(b);          // compile error: constexpr function call requires
constant expression argument
}

int main()
{
    constexpr int a { 5 };

    std::cout << foo(a); // okay: constant expression a can be used as argument to
constexpr function foo()

    return 0;
}
```

In the above example, function parameter `b` is not constexpr (even though argument `a` is a constant expression). This means `b` cannot be used anywhere a constant expression is required, such as the the initializer for a constexpr variable (e.g. `b2`) or in a call to a constexpr function (`goo(b)`).

Perhaps surprisingly, the parameters of a constexpr function are also not constexpr (even though constexpr functions can only be evaluated at compile-time). This decision was made for the sake of consistency.

The parameters of constexpr/constexpr functions may be declared as `const`, in which case they are treated as runtime constants.

Related content

If you need parameters that are constant expressions, see [11.9 -- Non-type template parameters](#).

Constexpr/constexpr functions can use non-const local variables

Within a constexpr or consteval function, we can use local variables that are not constexpr, and the value of these variables can be changed.

As a silly example:

```
#include <iostream>

consteval int doSomething(int x, int y) // function is consteval
{
    x = x + 2;           // we can modify the value of non-const function parameters

    int z { x + y }; // we can instantiate non-const local variables
    if (x > y)
        z = z - 1;    // and then modify their values

    return z;
}

int main()
{
    constexpr int g { doSomething(5, 6) };
    std::cout << g << '\n';

    return 0;
}
```

When such functions are evaluated at compile-time, the compiler will essentially “execute” the function and return the calculated value.

Constexpr/consteval functions can use function parameters and local variables as arguments in constexpr function calls

Above, we noted, “When a constexpr (or consteval) function is being evaluated at compile-time, any other functions it calls are required to be evaluated at compile-time.”

Perhaps surprisingly, a constexpr or consteval function can use its function parameters (which aren’t constexpr) or even local variables (which may not be const at all) as arguments in a constexpr function call. When a constexpr or consteval function is being evaluated at compile-time, the value of all function parameters and local variables must be known to the compiler (otherwise it couldn’t evaluate them at compile-time). Therefore, in this specific context, C++ allows these values to be used as arguments in a call to a constexpr function, and that constexpr function call can still be evaluated at compile-time.

```

#include <iostream>

constexpr int goo(int c) // goo() is now constexpr
{
    return c;
}

constexpr int foo(int b) // b is not a constant expression within foo()
{
    return goo(b);        // if foo() is resolved at compile-time, then `goo(b)` can
    also be resolved at compile-time
}

int main()
{
    std::cout << foo(5);

    return 0;
}

```

In the above example, `foo(5)` may or may not be evaluated at compile time. If it is, then the compiler knows that `b` is `5`. And even though `b` is not `constexpr`, the compiler can treat the call to `goo(b)` as if it were `goo(5)` and evaluate that function call at compile-time. If `foo(5)` is instead resolved at runtime, then `goo(b)` will also be resolved at runtime.

Can a `constexpr` function call a non-`constexpr` function?

The answer is yes, but only when the `constexpr` function is being evaluated in a non-constant context. A non-`constexpr` function may not be called when a `constexpr` function is evaluating in a constant context (because then the `constexpr` function wouldn't be able to produce a compile-time constant value), and doing so will produce a compilation error.

Calling a non-`constexpr` function is allowed so that a `constexpr` function can do something like this:

```

#include <type_traits> // for std::is_constant_evaluated

constexpr int someFunction()
{
    if (std::is_constant_evaluated()) // if compile-time evaluation
        return someConstexprFcn();  // calculate some value at compile time
    else                               // runtime evaluation
        return someNonConstexprFcn(); // calculate some value at runtime
}

```

Now consider this variant:

```
constexpr int someFunction(bool b)
{
    if (b)
        return someConstexprFcn();
    else
        return someNonConstexprFcn();
}
```

This is legal as long as `someFunction(false)` is never called in a constant expression.

As an aside...

Prior to C++23, the C++ standard says that a constexpr function must return a constexpr value for at least one set of arguments, otherwise it is technically ill-formed. Calling a non-constexpr function unconditionally in a constexpr function makes the constexpr function ill-formed. However, compilers are not required to generate errors or warnings for such cases -- therefore, the compiler probably won't complain unless you try to call such a constexpr function in a constant context. In C++23, this requirement was rescinded.

For best results, we'd advise the following:

1. Avoid calling non-constexpr functions from within a constexpr function if possible.
2. If your constexpr function requires different behavior for constant and non-constant contexts, conditionalize the behavior with `if (std::is_constant_evaluated())`.
3. Always test your constexpr functions in a constant context, as they may work when called in a non-constant context but fail in a constant context.

Why not constexpr every eligible function?

There are a few reasons you may not want to constexpr a function:

1. `constexpr` is part of the interface of a function. Once a function is made constexpr, it can be called by other constexpr functions or used in contexts that require constant expressions. Removing the `constexpr` later will break such code.
2. `constexpr` makes functions harder to debug because you can't inspect them at runtime.

But as a general rule, if you can `constexpr` a function, you should.

Best practice

Unless you have a specific reason not to, a function that can be made `constexpr` generally should be made `constexpr`.