

8.10 — For statements

 learncpp.com/cpp-tutorial/for-statements/

By far, the most utilized loop statement in C++ is the for-statement. The **for statement** (also called a **for loop**) is preferred when we have an obvious loop variable because it lets us easily and concisely define, initialize, test, and change the value of loop variables.

As of C++11, there are two different kinds of for-loops. We'll cover the classic for-statement in this lesson, and the newer range-based for-statement in a future lesson ([16.8 -- Range-based for loops \(for-each\)](#)) once we've covered some other prerequisite topics.

The for-statement looks pretty simple in abstract:

```
for (init-statement; condition; end-expression)
    statement;
```

The easiest way to initially understand how a for-statement works is to convert it into an equivalent while-statement:

```
{ // note the block here
    init-statement; // used to define variables used in the loop
    while (condition)
    {
        statement;
        end-expression; // used to modify the loop variable prior to reassessment of
the condition
    }
} // variables defined inside the loop go out of scope here
```

Evaluation of for-statements

A for-statement is evaluated in 3 parts:

First, the init-statement is executed. This only happens once when the loop is initiated. The init-statement is typically used for variable definition and initialization. These variables have “loop scope”, which really just is a form of block scope where these variables exist from the point of definition through the end of the loop statement. In our while-loop equivalent, you can see that the init-statement is inside a block that contains the loop, so the variables defined in the init-statement go out of scope when the block containing the loop ends.

Second, for each loop iteration, the condition is evaluated. If this evaluates to true, the statement is executed. If this evaluates to false, the loop terminates and execution continues with the next statement beyond the loop.

Finally, after the statement is executed, the end-expression is evaluated. Typically, this expression is used to increment or decrement the loop variables defined in the init-statement. After the end-expression has been evaluated, execution returns to the second step (and the condition is evaluated again).

Let's take a look at a sample for-loop and discuss how it works:

```
#include <iostream>

int main()
{
    for (int i{ 1 }; i <= 10; ++i)
        std::cout << i << ' ';

    std::cout << '\n';

    return 0;
}
```

First, we declare a loop variable named `i`, and initialize it with the value `1`.

Second, `i <= 10` is evaluated, and since `i` is `1`, this evaluates to `true`. Consequently, the statement executes, which prints `1` and a space.

Finally, `++i` is evaluated, which increments `i` to `2`. Then the loop goes back to the second step.

Now, `i <= 10` is evaluated again. Since `i` has value `2`, this evaluates to `true`, so the loop iterates again. The statement prints `2` and a space, and `i` is incremented to `3`. The loop continues to iterate until eventually `i` is incremented to `11`, at which point `i <= 10` evaluates to `false`, and the loop exits.

Consequently, this program prints the result:

```
1 2 3 4 5 6 7 8 9 10
```

For the sake of example, let's convert the above for-loop into an equivalent while-loop:

```

#include <iostream>

int main()
{
    { // the block here ensures block scope for i
        int i{ 1 }; // our init-statement
        while (i <= 10) // our condition
        {
            std::cout << i << ' '; // our statement
            ++i; // our end-expression
        }
    }

    std::cout << '\n';
}

```

That doesn't look so bad, does it? Note that the outer braces are necessary here, because `i` goes out of scope when the loop ends.

For-loops can be hard for new programmers to read -- however, experienced programmers love them because they are a very compact way to do loops with a counter, with all of the necessary information about the loop variables, loop conditions, and loop variable modifiers presented up front. This helps reduce errors.

More for-loop examples

Here's an example of a function that uses a for-loop to calculate integer exponents:

```

#include <cstdint> // for fixed-width integers

// returns the value base ^ exponent -- watch out for overflow!
std::int64_t pow(int base, int exponent)
{
    std::int64_t total{ 1 };

    for (int i{ 0 }; i < exponent; ++i)
        total *= base;

    return total;
}

```

This function returns the value $\text{base}^{\text{exponent}}$ (base to the exponent power).

This is a straightforward incrementing for-loop, with `i` looping from `0` up to (but excluding) `exponent`.

If `exponent` is `0`, the for-loop will execute `0` times, and the function will return `1`.

If `exponent` is `1`, the for-loop will execute `1` time, and the function will return `1 * base`.

If `exponent` is `2`, the for-loop will execute `2` times, and the function will return `1 * base * base`.

Although most for-loops increment the loop variable by 1, we can decrement it as well:

```
#include <iostream>

int main()
{
    for (int i{ 9 }; i >= 0; --i)
        std::cout << i << ' ';

    std::cout << '\n';

    return 0;
}
```

This prints the result:

9 8 7 6 5 4 3 2 1 0

Alternately, we can change the value of our loop variable by more than 1 with each iteration:

```
#include <iostream>

int main()
{
    for (int i{ 0 }; i <= 10; i += 2) // increment by 2 each iteration
        std::cout << i << ' ';

    std::cout << '\n';

    return 0;
}
```

This prints the result:

0 2 4 6 8 10

The perils of **operator !=** in for-loop conditions

When writing a loop condition involving a value, we can often write the condition in many different ways. The following two loops execute identically:

```
#include <iostream>

int main()
{
    for (int i { 0 }; i < 10; ++i) // uses <
        std::cout << i;

    for (int i { 0 }; i != 10; ++i) // uses !=
        std::cout << i;

    return 0;
}
```

So which should we prefer? The former is the better choice, as it will terminate even if **i** jumps over the value **10**, whereas the latter will not. The following example demonstrates this:

```
#include <iostream>

int main()
{
    for (int i { 0 }; i < 10; ++i) // uses <, still terminates
    {
        std::cout << i;
        if (i == 9) ++i; // jump over value 10
    }

    for (int i { 0 }; i != 10; ++i) // uses !=, infinite loop
    {
        std::cout << i;
        if (i == 9) ++i; // jump over value 10
    }

    return 0;
}
```

Best practice

Avoid **operator !=** when doing numeric comparisons in the for-loop condition. Prefer **operator <** or **operator <=** where possible.

Off-by-one errors

One of the biggest problems that new programmers have with for-loops (and other loops that utilize counters) are off-by-one errors. **Off-by-one errors** occur when the loop iterates one too many or one too few times to produce the desired result.

Here's an example:

```

#include <iostream>

int main()
{
    // oops, we used operator< instead of operator<=
    for (int i{ 1 }; i < 5; ++i)
    {
        std::cout << i << ' ';
    }

    std::cout << '\n';

    return 0;
}

```

This program is supposed to print **1 2 3 4 5**, but it only prints **1 2 3 4** because we used the wrong relational operator.

Although the most common cause for these errors is using the wrong relational operator, they can sometimes occur by using pre-increment or pre-decrement instead of post-increment or post-decrement, or vice-versa.

Omitted expressions

It is possible to write *for loops* that omit any or all of the statements or expressions. For example, in the following example, we'll omit the init-statement and end-expression, leaving only the condition:

```

#include <iostream>

int main()
{
    int i{ 0 };
    for ( ; i < 10; ) // no init-statement or end-expression
    {
        std::cout << i << ' ';
        ++i;
    }

    std::cout << '\n';

    return 0;
}

```

This *for loop* produces the result:

```
0 1 2 3 4 5 6 7 8 9
```

Rather than having the *for loop* do the initialization and incrementing, we've done it manually. We have done so purely for academic purposes in this example, but there are cases where not defining a loop variable (because you already have one) or not incrementing it in the end-expression (because you're incrementing it some other way) is desired.

Although you do not see it very often, it is worth noting that the following example produces an infinite loop:

```
for (;;)
    statement;
```

The above example is equivalent to:

```
while (true)
    statement;
```

This might be a little unexpected, as you'd probably expect an omitted condition-expression to be treated as **false**. However, the C++ standard explicitly (and inconsistently) defines that an omitted condition-expression in a for-loop should be treated as **true**.

We recommend avoiding this form of the for loop altogether and using **while (true)** instead.

For-loops with multiple counters

Although for-loops typically iterate over only one variable, sometimes for-loops need to work with multiple variables. To assist with this, the programmer can define multiple variables in the init-statement, and can make use of the comma operator to change the value of multiple variables in the end-expression:

```
#include <iostream>

int main()
{
    for (int x{ 0 }, y{ 9 }; x < 10; ++x, --y)
        std::cout << x << ' ' << y << '\n';

    return 0;
}
```

This loop defines and initializes two new variables: **x** and **y**. It iterates **x** over the range **0** to **9**, and after each iteration **x** is incremented and **y** is decremented.

This program produces the result:

```
0 9
1 8
2 7
3 6
4 5
5 4
6 3
7 2
8 1
9 0
```

This is about the only place in C++ where defining multiple variables in the same statement, and use of the comma operator is considered an acceptable practice.

Related content

We cover the comma operator in lesson [6.5 -- The comma operator](#).

Best practice

Defining multiple variables (in the init-statement) and using the comma operator (in the end-expression) is acceptable inside a for-statement.

Nested for-loops

Like other types of loops, for-loops can be nested inside other loops. In the following example, we're nesting a for-loop inside another for-loop:

```
#include <iostream>

int main()
{
    for (char c{ 'a' }; c <= 'e'; ++c) // outer loop on letters
    {
        std::cout << c; // print our letter first

        for (int i{ 0 }; i < 3; ++i) // inner loop on all numbers
            std::cout << i;

        std::cout << '\n';
    }

    return 0;
}
```

For each iteration of the outer loop, the inner loop runs in its entirety. Consequently, the output is:


```
a012
b012
c012
d012
e012
```

Here's some more detail on what's happening here. The outer loop runs first, and char `c` is initialized to `'a'`. Then `c <= 'e'` is evaluated, which is `true`, so the loop body executes. Since `c` is set to `'a'`, this first prints `a`. Next the inner loop executes entirely (which prints `0`, `1`, and `2`). Then a newline is printed. Now the outer loop body is finished, so the outer loop returns to the top, `c` is incremented to `'b'`, and the loop condition is re-evaluated. Since the loop condition is still `true` the next iteration of the outer loop begins. This prints `b012\n`. And so on.

Variables used only inside a loop should be defined inside the loop.

New programmers often think that creating variables is expensive, so that it is better to create a variable once (and then assign values to it) than create it many times (and use initialization). That leads to loops that look like some variation of the following:

```
#include <iostream>

int main()
{
    int i {}; // i defined outside loop
    for (i = 0; i < 10; ++i) // i assigned value
    {
        std::cout << i << ' ';
    }

    // i can still be accessed here

    std::cout << '\n';

    return 0;
}
```

However, creating a variable has no cost -- it is the initialization that has a cost, and there is typically no cost difference between initialization and assignment. The above example makes `i` available beyond the loop. Unless using a variable beyond the loop is required, defining a variable outside the loop is likely to have two consequences:

1. It makes our program more complex, as we have to read more code to see where the variable is used.
2. It may actually be slower because the compiler may not be able to optimize a variable with a larger scope as effectively.

Consistent with our best practice to define variables in the smallest reasonable scope possible, a variable that is only used within a loop should be defined inside the loop rather than outside.

Best practice

Variables used only inside a loop should be defined inside the scope of the loop.

Conclusion

For-statements are the most commonly used loop in the C++ language. Even though its syntax is typically a bit confusing to new programmers, you will see for-loops so often that you will understand them in no time at all!

For-statements excel when you have a counter variable. If you do not have a counter, a while-statement is probably a better choice.

Best practice

Prefer for-loops over while-loops when there is an obvious loop variable.
Prefer while-loops over for-loops when there is no obvious loop variable.

Quiz time

Question #1

Write a for-loop that prints every even number from 0 to 20.

Show Solution

Question #2

Write a function named `sumTo()` that takes an integer parameter named `value`, and returns the sum of all the numbers from 1 to `value`.

For example, `sumTo(5)` should return 15, which is $1 + 2 + 3 + 4 + 5$.

Hint: Use a non-loop variable to accumulate the sum as you iterate from 1 to the input value, much like the `pow()` example above uses the `total` variable to accumulate the return value each iteration.

Show Solution

Question #3

What's wrong with the following for loop?

```
// Print all numbers from 9 to 0
for (unsigned int i{ 9 }; i >= 0; --i)
    std::cout << i<< ' ';
```

Show Solution

Question #4

Fizz Buzz is a simple math game used to teach children about divisibility. It is also sometimes used as an interview question to assess basic programming skills.

The rules of the game are simple: Starting at 1, and counting upward, replace any number divisible only by three with the word “fizz”, any number only divisible by five with the word “buzz”, and any number divisible by both 3 and 5 with the word “fizzbuzz”.

Implement this game inside a function named `fizzbuzz()` that takes a parameter determining what number to count up to. Use a for-loop and a single if-else chain (meaning you can use as many else-if as you like).

The output of `fizzbuzz(15)` should match the following:

```
1
2
fizz
4
buzz
fizz
7
8
fizz
buzz
11
fizz
13
14
fizzbuzz
```

Show Solution

Question #5

Modify the FizzBuzz program you wrote in the previous quiz to add the rule that numbers divisible by seven should be replaced with “pop”. Run the program for 150 iterations.

In this version, using an if/else chain to explicitly cover every possible combination of words will result in a function that is too long. Optimize your function so only 4 if-statements are used: one for each of the non-compound words (“fizz”, “buzz”, “pop”), and one for the case where a number is printed.

Show Hint

Here are some snippets of the expected output:

```
4  
buzz  
fizz  
pop  
8
```

```
19  
buzz  
fizzpop  
22
```

```
104  
fizzbuzzpop  
106
```

Show Solution